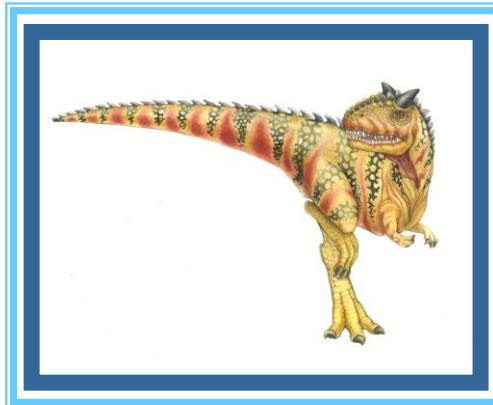


Chapter 3: Processes





Process Concept

- An operating system executes a variety of programs:
 - Batch system – **jobs**
 - Time-shared systems – **user programs** or **tasks**
- Textbook uses the terms **job** and **process** almost interchangeably
- **Process** – a **program in execution**; process execution must progress in **sequential** fashion
- Multiple parts
 - The **program code**, also called **text section**
 - Current activity including **program counter**, processor registers
 - **Stack** containing **temporary data**
 - Function **parameters**, **return addresses**, **local variables**
 - **Data section** containing **global variables**
 - **Heap** containing memory **dynamically allocated** during run time





Process Concept (Cont.)

- Program is **passive** entity stored on disk (**executable file**), process is **active**
 - Program becomes **process** when **executable file loaded into memory**
- Execution of program started via GUI mouse clicks, command line entry of its name, etc
- One program can be several processes
 - Consider **multiple users executing the same program**

In a Bach system, jobs are processed one after the other in a batch mode, where each job is executed to completion before the next one starts.

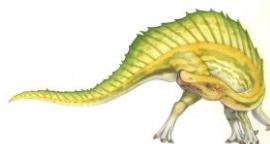

In a time-sharing system, multiple users can interact with the computer simultaneously, and the system rapidly switches between different user tasks, giving each user a small time slice (time quantum) for execution.

Time-sharing systems are designed to provide the illusion of concurrent execution for multiple users, even though the computer's central processing unit (CPU) can only execute one instruction at a time.






Process in Memory





Process in Memory






Process State

- As a process executes, it changes **state**
 - **new**: The process is being created
 - **running**: Instructions are being executed
 - **waiting**: The process is waiting for some event to occur
 - **ready**: The process is waiting to be assigned to a processor
 - **terminated**: The process has finished execution





Diagram of Process State






Process Control Block (PCB)


Information associated with each process
(also called **task control block**)


- Process state – running, waiting, etc
- Program counter – location of instruction to next execute
- CPU registers – contents of all process-centric registers
- CPU scheduling information- priorities, scheduling queue pointers
- Memory-management information – memory allocated to the process
- Accounting information – CPU used, clock time elapsed since start, time limits
- I/O status information – I/O devices allocated to process, list of open files





CPU Switch From Process to Process





Threads

- So far, process has a single thread of execution
- Consider having **multiple program counters per process**
 - Multiple locations can execute at once
 - ▶ Multiple threads of control -> **threads**
- Must then have storage for thread details, multiple program counters in PCB



A process is an independent program or unit of execution that consists of its own memory space, resources, and system information. Processes are isolated from each other and run independently, which means they don't share memory directly.

A thread is a lightweight unit of execution that exists within a process. Threads share the same memory space and resources as the parent process. They are used to perform separate tasks concurrently within a single program



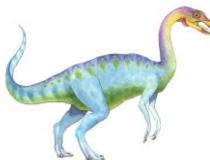


Threads



threaded

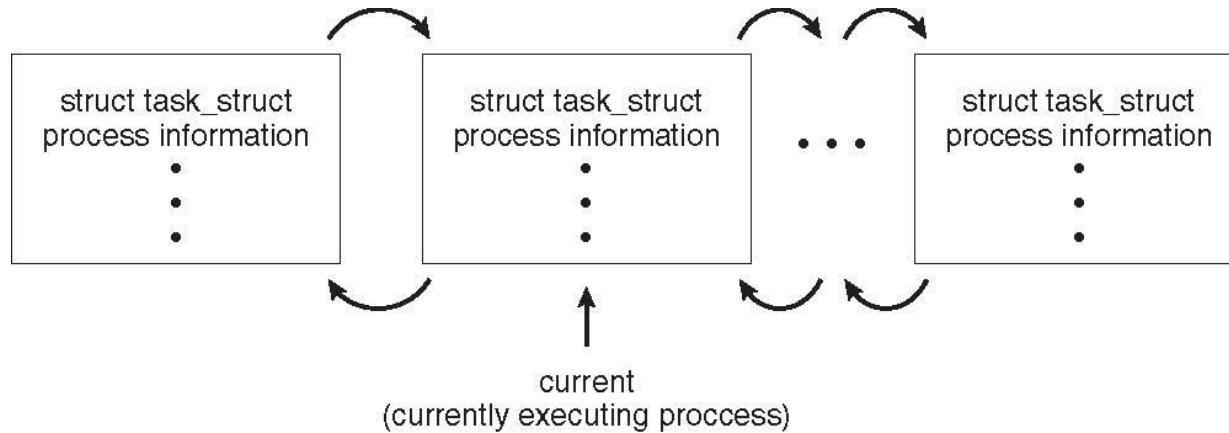


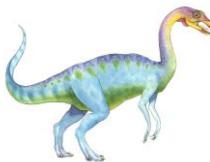


Process Representation in Linux

Represented by the C structure task_struct

```
pid t_pid; /* process identifier */  
long state; /* state of the process */  
unsigned int time_slice /* scheduling information */  
struct task_struct *parent; /* this process's parent */  
struct list_head children; /* this process's children */  
struct files_struct *files; /* list of open files */  
struct mm_struct *mm; /* address space of this process */
```





Process Scheduling

- Maximize CPU use, quickly switch processes onto CPU for time sharing
- **Process scheduler** selects among available processes for next execution on CPU
- Maintains **scheduling queues** of processes
 - **Job queue** – set of all processes in the system
 - **Ready queue** – set of all processes residing in main memory, ready and waiting to execute
 - **Device queues** – set of processes waiting for an I/O device

Key Differences: Processes migrate among the various queues

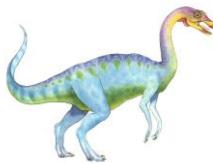
Isolation: Processes are isolated from each other, while threads within the same process share the same memory space and resources.

Resource Overhead: Processes have a higher resource overhead because they require separate memory and resources. Threads have lower overhead since they share the parent process's resources.

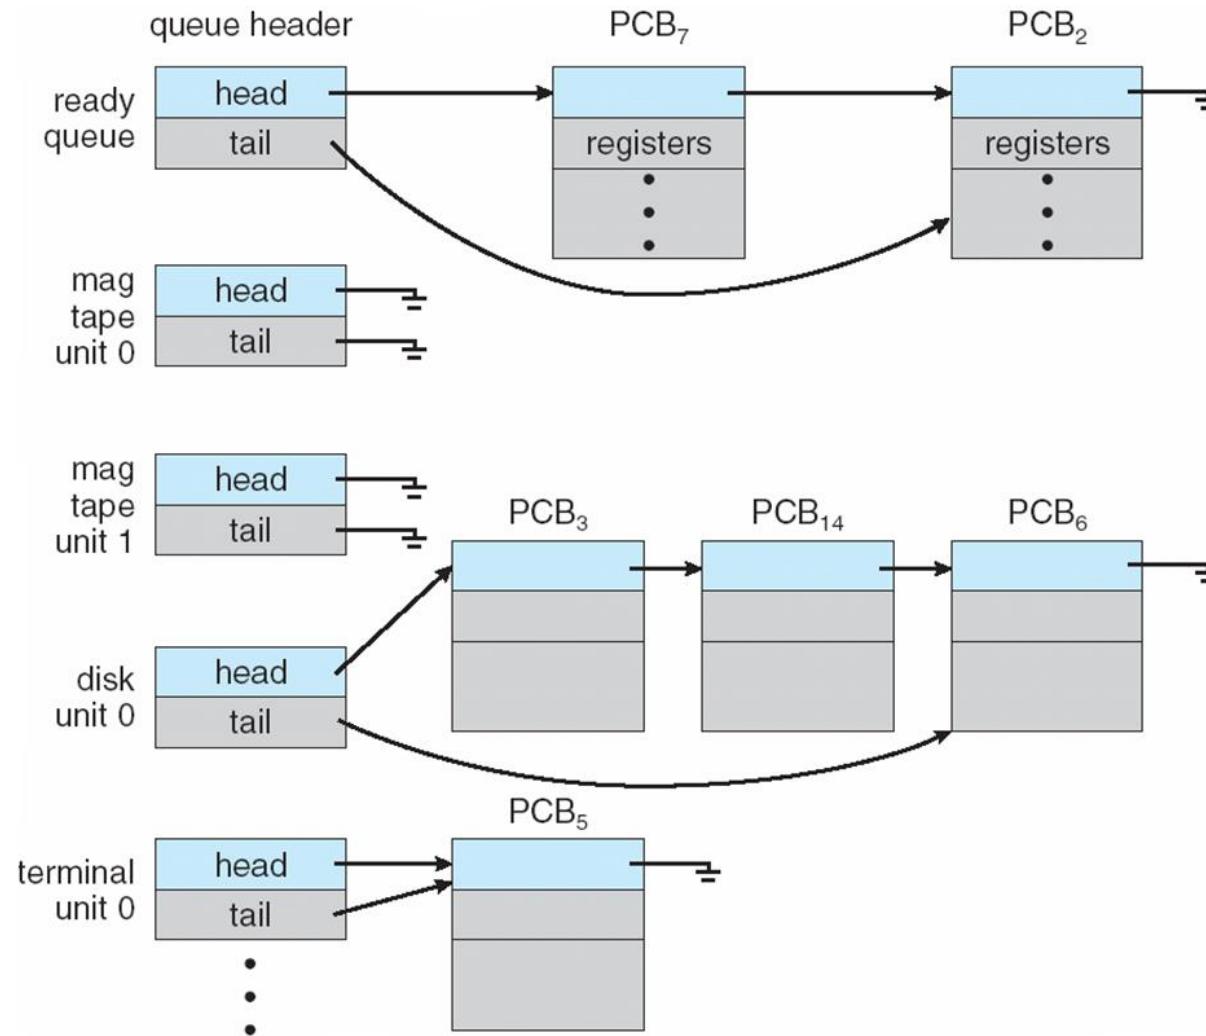
Communication: Communication between processes often requires more complex mechanisms like IPC, whereas communication between threads can be achieved through shared memory, making it more straightforward.

Fault Tolerance: If one process crashes, it generally does not affect other processes. In contrast, if one thread in a process crashes, it can potentially impact other threads in the same process.

Concurrency: Threads within a single process can achieve true parallelism on multi-core processors, as they can run in parallel. Processes can run concurrently but not necessarily in parallel, as they may have to share CPU time.




Ready Queue And Various I/O Device Queues





Representation of Process Scheduling

- Queueing diagram represents queues, resources, flows





Schedulers


- **Short-term scheduler** (or **CPU scheduler**) – selects which process should be executed next and allocates CPU
 - Sometimes the only scheduler in a system
 - Short-term scheduler is invoked frequently (milliseconds) ⇒ (must be fast)
- **Long-term scheduler** (or **job scheduler**) – selects which processes should be brought into the ready queue
 - Long-term scheduler is invoked infrequently (seconds, minutes) ⇒ (may be slow)
 - The long-term scheduler controls the **degree of multiprogramming**
- Processes can be described as either:
 - **I/O-bound process** – spends more time doing I/O than computations, many short CPU bursts
 - **CPU-bound process** – spends more time doing computations; few very long CPU bursts
- Long-term scheduler strives for good **process mix**






Addition of Medium Term Scheduling

- **Medium-term scheduler** can be added if **degree of multiple programming** needs to **decrease**
 - Remove process from **memory**, store on **disk**, bring back in from **disk** to continue execution: **swapping**





Addition of Medium Term Scheduling





Multitasking in Mobile Systems

- Some mobile systems (e.g., early version of iOS) allow only one process to run, others suspended
- Due to screen real estate, user interface limits, iOS provides for a
 - Single **foreground** process- controlled via user interface
 - Multiple **background** processes– in memory, running, but not on the display, and with limits
 - Limits include single, short task, receiving notification of events, specific long-running tasks like audio playback
- Android runs foreground and background, with fewer limits
 - Background process uses a **service** to perform tasks
 - Service can keep running even if background process is suspended
 - Service has no user interface, small memory use





Context Switch

- When CPU switches to another process, the system must **save the state** of the old process and load the **saved state** for the new process via a **context switch**
- **Context** of a process represented in the PCB
- Context-switch time is **overhead**; the system does no useful work while switching
 - The more **complex** the OS and the PCB → the **longer** the context **switch**
- Time dependent on **hardware support**
 - Some hardware provides **multiple sets of registers per CPU** → **multiple contexts loaded at once**





Operations on Processes

- System must provide mechanisms for:
 - process creation,
 - process termination,
 - and so on as detailed next

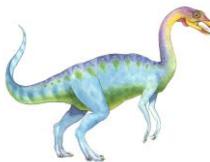




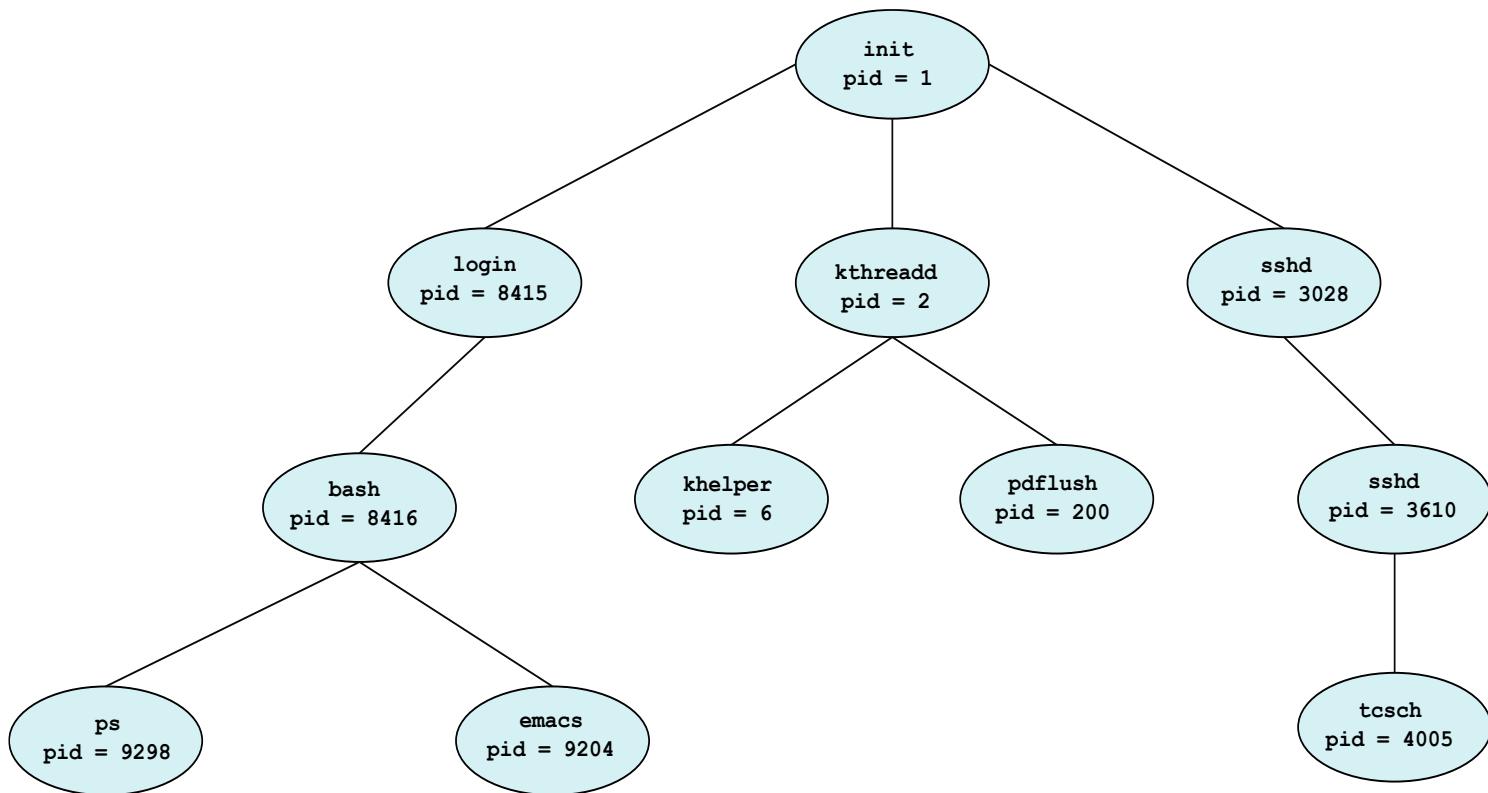
Process Creation

- Parent process create children processes, which, in turn create other processes, forming a tree of processes
- Generally, process identified and managed via a process identifier (pid)
- Resource sharing options
 - Parent and children share all resources
 - Children share subset of parent's resources
 - Parent and child share no resources
- Execution options
 - Parent and children execute concurrently
 - Parent waits until children terminate






A Tree of Processes in Linux





Process Creation (Cont.)

- Address space options
 - Child duplicate of parent
 - Child has a program loaded into it
- UNIX examples
 - `fork()` system call creates new process
 - `exec()` system call used after a `fork()` to replace the process' memory space with a new program





C Program Forking Separate Process

```
#include <sys/types.h>
#include <stdio.h>
#include <unistd.h>

int main()
{
    pid_t pid;

    /* fork a child process */
    pid = fork();

    if (pid < 0) { /* error occurred */
        fprintf(stderr, "Fork Failed");
        return 1;
    }
    else if (pid == 0) { /* child process */
        execlp("/bin/ls", "ls", NULL);
    }
    else { /* parent process */
        /* parent will wait for the child to complete */
        wait(NULL);
        printf("Child Complete");
    }
}

return 0;
}
```





Process Termination

- Process executes last statement and then asks the operating system to delete it using the `exit()` system call.
 - Returns status data from child to parent (via `wait()`)
 - Process' resources are deallocated by operating system
- Parent may terminate the execution of children processes using the `abort()` system call. Some reasons for doing so:
 - Child has exceeded allocated resources
 - Task assigned to child is no longer required
 - The parent is exiting and the operating systems does not allow a child to continue if its parent terminates





Process Termination

- Some operating systems do not allow child to exists if its parent has terminated. If a process terminates, then all its children must also be terminated.
 - **cascading termination.** All children, grandchildren, etc. are terminated.
 - The termination is initiated by the operating system.
- The parent process may **wait for termination of a child** process by using the **wait()** system call. The call returns status information and the pid of the terminated process

```
pid = wait(&status);
```
- If **no parent waiting** (did not invoke **wait()**) process is a **zombie**
- If **parent terminated without invoking wait**, process is an **orphan**





Multiprocess Architecture – Chrome Browser

- Many web browsers ran as single process (some still do)
 - If one web site causes trouble, entire browser can hang or crash
- Google Chrome Browser is multiprocess with 3 different types of processes:
 - **Browser** process manages user interface, disk and network I/O
 - **Renderer** process renders web pages, deals with HTML, Javascript. A new renderer created for each website opened
 - ▶ Runs in **sandbox** restricting disk and network I/O, minimizing effect of security exploits
 - **Plug-in** process for each type of plug-in





Interprocess Communication


- Processes within a system may be **independent** or **cooperating**
- **Cooperating** process can affect or be affected by other processes, including sharing data
- **Independent** process cannot affect or be affected by the execution of another process
- Reasons for cooperating processes:
 - Information sharing
 - Computation speedup
 - Modularity
- Cooperating processes need **interprocess communication (IPC)**
- Two models of IPC
 - **Shared memory**
 - **Message passing**





Communications Models

(a) Message passing. (b) shared memory.





Producer-Consumer Problem

- Paradigm for cooperating processes, *producer* process produces information that is consumed by a *consumer* process
 - **unbounded-buffer** places no practical limit on the size of the buffer
 - **bounded-buffer** assumes that there is a fixed buffer size





Bounded-Buffer – Shared-Memory Solution

- Shared data

```
#define BUFFER_SIZE 10

typedef struct {

    . . .

} item;

item buffer[BUFFER_SIZE];
int in = 0;    => next Empty slot in Buffer
int out = 0;   => next item in buffer
```

- Solution is correct, but can only use BUFFER_SIZE-1 elements





Bounded-Buffer – Producer

```
item next_produced;  
while (true) {  
    /* produce an item in next produced */  
    while (((in + 1) % BUFFER_SIZE) == out)  
if Buffer Full => wait in while loop => can't insert an item in buffer  
        ; /* do nothing */  
    buffer[in] = next_produced;           Buffer not Full => insert  
    in = (in + 1) % BUFFER_SIZE;         NextProduced item in buffer.  
}
```

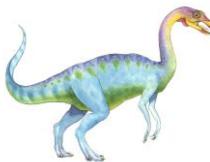




Bounded Buffer – Consumer

```
item next_consumed;  
while (true) {  
    while (in == out)      if buffer is Empty => can't Remove any item => wait  
                           in while loop  
        ; /* do nothing */  
    next_consumed = buffer[out];      Buffer not Empty => remove An  
    out = (out + 1) % BUFFER_SIZE;   item from buffer.  
  
    /* consume the item in next consumed */  
}
```





Interprocess Communication – Shared Memory

- An area of memory shared among the processes that wish to communicate
- The communication is under the control of the users processes not the operating system.
- Major issues is to provide mechanism that will allow the user processes to synchronize their actions when they access shared memory.
- Synchronization is discussed in great details in Chapter 5.





Interprocess Communication – Message Passing

- Mechanism for processes to communicate and to synchronize their actions
- Message system – processes communicate with each other without resorting to shared variables
- IPC facility provides two operations:
 - `send(message)`
 - `receive(message)`
- The message size is either fixed or variable





Message Passing (Cont.)

- If processes P and Q wish to communicate, they need to:
 - Establish a **communication link** between them
 - Exchange messages via send/receive
- Implementation issues:
 - How are **links established**?
 - Can a link be **associated** with **more** than two processes?
 - **How many links** can there be between **every pair of** communicating processes?
 - What is the **capacity of a link**?
 - Is the **size of a message** that the link can **accommodate** fixed or variable?
 - Is a link **unidirectional** or **bi-directional**?





Message Passing (Cont.)

- Implementation of communication link
 - Physical:
 - ▶ Shared memory
 - ▶ Hardware bus
 - ▶ Network
 - Logical:
 - ▶ Direct or indirect
 - ▶ Synchronous or asynchronous
 - ▶ Automatic or explicit buffering

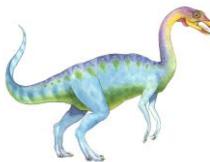




Direct Communication

- Processes must name each other explicitly:
 - **send** (P , message) – send a message to process P
 - **receive**(Q , message) – receive a message from process Q
- Properties of communication link
 - Links are established automatically
 - A link is associated with exactly one pair of communicating processes
 - Between each pair there exists exactly one link
 - The link may be unidirectional, but is usually bi-directional





Indirect Communication

- Messages are directed and received from mailboxes (also referred to as ports)
 - Each mailbox has a unique id
 - Processes can communicate only if they share a mailbox
- Properties of communication link
 - Link established only if processes share a common mailbox
 - A link may be associated with many processes
 - Each pair of processes may share several communication links
 - Link may be unidirectional or bi-directional





Indirect Communication

- Operations

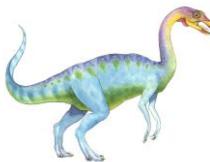
- **create** a **new** mailbox (port)
 - **send and receive** messages through mailbox
 - **destroy** a mailbox

- Primitives are defined as:

send(*A, message*) – send a message **to** mailbox *A*

receive(*A, message*) – receive a message **from** mailbox *A*





Indirect Communication

- Mailbox sharing
 - P_1 , P_2 , and P_3 share mailbox A
 - P_1 , sends; P_2 and P_3 receive
 - Who gets the message?
- Solutions
 - Allow a link to be associated with at most two processes
 - Allow only one process at a time to execute a receive operation
 - Allow the system to select arbitrarily the receiver.
Sender is notified who the receiver was.

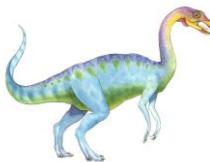




Synchronization

- ❑ Message passing may be either **blocking** or **non-blocking**
- ❑ **Blocking** is considered **synchronous**
 - ❑ **Blocking send** -- the **sender** is blocked until the message is **received**
 - ❑ **Blocking receive** -- the **receiver** is blocked until a **message** is **available**
- ❑ **Non-blocking** is considered **asynchronous**
 - ❑ **Non-blocking send** -- the **sender** sends the message and **continues**
 - ❑ **Non-blocking receive** -- the **receiver** receives:
 - ❑ A **valid message**, or
 - ❑ **Null message**
- ❑ Different combinations possible
 - ❑ If both **send** and **receive** are **blocking**, we have a **rendezvous**





Synchronization (Cont.)

- Producer-consumer becomes trivial

```
message next_produced;  
while (true) {  
    /* produce an item in next produced */  
    send(next_produced);  
}
```

```
message next_consumed;  
while (true) {  
    receive(next_consumed);  
  
    /* consume the item in next consumed */  
}
```





Buffering

- Queue of messages attached to the link.
- implemented in one of three ways
 1. Zero capacity – no messages are queued on a link.
Sender must wait for receiver (rendezvous)
 2. Bounded capacity – finite length of n messages
Sender must wait if link full
 3. Unbounded capacity – infinite length
Sender never waits





Communications in Client-Server Systems

- Sockets
- Remote Procedure Calls
- Pipes
- Remote Method Invocation (Java)






Sockets

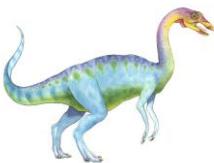
- A **socket** is defined as an endpoint for communication
- Concatenation of IP address and **port** – a number included at start of message packet to differentiate network services on a host
- The socket **161.25.19.8:1625** refers to port **1625** on host **161.25.19.8**
- Communication consists between a pair of sockets
- All ports below 1024 are **well known**, used for standard services
- Special IP address 127.0.0.1 (**loopback**) to refer to system on which process is running





Socket Communication





Pipes

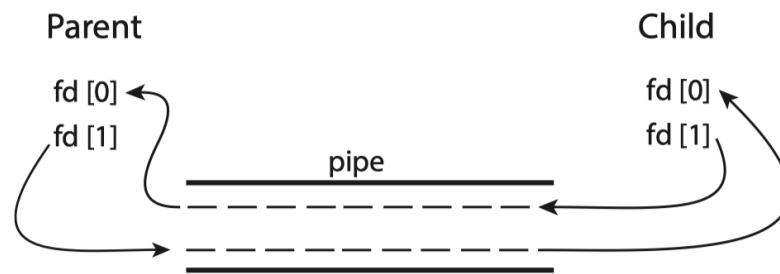
- Acts as a conduit allowing two processes to communicate
- Issues:
 - Is communication unidirectional or bidirectional?
 - In the case of two-way communication, is it half or full-duplex?
 - Must there exist a relationship (i.e., **parent-child**) between the communicating processes?
 - Can the pipes be used over a network?
- Ordinary pipes – cannot be accessed from outside the process that created it. Typically, a parent process creates a pipe and uses it to communicate with a child process that it created.
- Named pipes – can be accessed without a parent-child relationship.





Ordinary Pipes

- ❑ Ordinary Pipes allow communication in standard producer-consumer style
- ❑ Producer writes to one end (the **write-end** of the pipe)
- ❑ Consumer reads from the other end (the **read-end** of the pipe)
- ❑ Ordinary pipes are therefore unidirectional
- ❑ Require parent-child relationship between communicating processes



- ❑ Windows calls these **anonymous pipes**
- ❑ See Unix and Windows code samples in textbook





Named Pipes

- Named Pipes are more powerful than ordinary pipes
- Communication is bidirectional
- No parent-child relationship is necessary between the communicating processes
- Several processes can use the named pipe for communication
- Provided on both UNIX and Windows systems



End of Chapter 3

